

Kubernetes	2
Lab 1 Installing Kubernetes	3
Lab 2 Kubernetes Cluster Details	6
Lab 3 Namespaces	7
Lab 4 Create PODs using RUN command	8
Lab 5 Creating PODs using POD Object	9
Lab 6 Create PODs using ReplicationSet	10
Lab 7 Create PODs using DeploymentSet	11
Lab 8 Managing PODs	13
Lab 9 Exposing PODs to Internet	16
Lab 10 Kubernetes Networking	22
Lab 12 Taints and Tolerations	24
Lab 15 Kubernetes Backup	27

Lab 1 Installing Kubernetes

Pre-requisites

- Create 3 VMs in the same Azure Region
- Size: DC1s_v3
- Location: UK South (or any)
- Ubuntu 22.04 LTS
- Username/Password
- VM names: master, node1, node2
- VNET: Accept the default 10.0.0.0/16

Install Kubernetes on master, node and node2

Execute these commands on all 3 vm's

```
1 # Update the VMs
2 sudo apt update
3 sudo apt-upgrade -y
4 sudo apt-get remove needrestart -y
5
6 # Add entries in hosts file
7 sudo sed -i '2i10.0.0.4 master' /etc/hosts
8 sudo sed -i '3i10.0.0.5 node1' /etc/hosts
9 sudo sed -i '4i10.0.0.6 node2' /etc/hosts
10
11 # Turn off swap (Most cloud VMs don't have swap turned on by default, but still)
12 sudo swapoff -a
13
14 # Adjust kernel settings
15 sudo tee /etc/modules-load.d/containerd.conf <<EOF
16 overlay
17 br_netfilter
18 EOF
19
20 sudo modprobe overlay
21 sudo modprobe br_netfilter
22
23 sudo tee /etc/sysctl.d/kubernetes.conf <<EOF
24 net.bridge.bridge-nf-call-ip6tables = 1
25 net.bridge.bridge-nf-call-iptables = 1
26 net.ipv4.ip_forward = 1
27 EOF
28
29 # Reload system
30 sudo sysctl --system
31
32 # Install a few utils
33 sudo apt install -y curl gnupg2 software-properties-common apt-transport-https ca-certificates
34
35 sudo curl -fsSL https://download.docker.com/linux/ubuntu/gpg \
36 | sudo gpg --dearmor -o /etc/apt/trusted.gpg.d/docker.gpg
37
```

```

38 sudo add-apt-repository "deb [arch=amd64] https://download.docker.com/linux/ubuntu $(lsb_release -cs) stable"
39 -y
40 # Update the system
41 sudo apt update
42
43 # Install container runtime
44 sudo apt install -y containerd.io
45
46 containerd config default | sudo tee /etc/containerd/config.toml >/dev/null 2>&1
47
48 sudo sed -i 's/SystemdCgroup \= false/SystemdCgroup \= true/g' /etc/containerd/config.toml
49 sudo systemctl restart containerd
50 sudo systemctl enable containerd
51
52 # Install kubernetes
53 sudo apt-get update
54
55 curl -fsSL https://pkgs.k8s.io/core:/stable:/v1.29/deb/Release.key \
56 | sudo gpg --dearmor -o /etc/apt/keyrings/kubernetes-apt-keyring.gpg
57
58 echo 'deb [signed-by=/etc/apt/keyrings/kubernetes-apt-keyring.gpg] \
59 https://pkgs.k8s.io/core:/stable:/v1.29/deb/ /' | sudo tee /etc/apt/sources.list.d/kubernetes.list
60
61 sudo apt-get update
62 sudo apt-get install -y kubelet kubeadm kubectl
63 sudo apt-mark hold kubelet kubeadm kubectl

```

Now create the cluster

Perform only on master node

```

1 # Configure Kubernetes on Master Node
2 # PERFORM THIS ON CONTROL-PLANE OR MASTER NODES ONLY
3 sudo kubeadm init --control-plane-endpoint=master
4
5 # If you have only 1 CPU for master node, use this command
6 sudo kubeadm init --control-plane-endpoint=master --ignore-preflight-errors=NumCPU
7
8 mkdir -p $HOME/.kube
9 sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
10 sudo chown $(id -u):$(id -g) $HOME/.kube/config
11
12 # Install Calico Pod Network Add-on
13 kubectl apply -f https://raw.githubusercontent.com/projectcalico/calico/v3.25.0/manifests/calico.yaml

```

Print the join command on the master node

Perform only on master node

```

1 # List existing tokens (optional)
2 kubeadm token list
3
4 # To print a join command for a new worker node use
5 kubeadm token create --print-join-command

```

Join nodes to cluster

Perform on both node and node2

Paste the output of earlier step and it looks like this.

```
1 # Join worker node to kubernetes cluster
2 # On the master node copy these lines and paste here. Your master node may generate a different output.
3 sudo kubeadm join master:6443 --token uf71xc.3g7yt5vl9zwdqc1a \
4     --discovery-token-ca-cert-hash sha256:9aaba87d3dd19be12234eda79075a76bb6ebb6cdab1e0442c9bba2b33b63fbd7
5
6
```

Verify cluster

```
1 kubectl get nodes
2 kubectl get pods -n kube-system
```

Lab 2 Kubernetes Cluster Details

Get cluster details

```
1 kubectl cluster-info
```

List nodes

```
1 kubectl get nodes
```

List namespaces in the cluster

```
1 kubectl get namespaces
```

Get current deployments in the cluster

```
1 kubectl get deployments --all-namespaces
```

Get services in the cluster

```
1 kubectl get services --all-namespaces
```

Lab 3 Namespaces

List existing namespaces

```
1 kubectl get namespaces
```

Create a New Namespace

Find the API version to use for namespace

```
1 kubectl api-resources | grep namespaces
```

Create a YAML file called as `dev.yaml`

```
1 nano dev.yaml
```

```
1 apiVersion: v1
2 kind: Namespace
3 metadata:
4   name: dev
```

```
1 kubectl apply -f dev.yaml
```

Verify creation

```
1 kubectl get namespaces
```

Create namespace using create command

```
1 kubectl create namespace prod
```

Verify creation

```
1 kubectl get namespaces
```

Using a Github Repo link

```
1 kubectl apply -f \
2 https://raw.githubusercontent.com/reposforlabs/kubernetes/main/namespace.yaml
```

Verify creation

```
1 kubectl get ns
```

List pods in all namespaces

```
1 kubectl get pods --all-namespaces
```

Lab 4 Create PODs using RUN command

Using RUN command

In default namespace

```
1 kubectl run pod1 --image=nginx --port=80 --labels=owner=raghavendra
```

List

```
1 kubectl get pods
```

In dev namespace

```
1 kubectl run pod1 --image=nginx --port=80 --labels=type=dev -n dev
```

List

```
1 kubectl get pods -n dev
```

Find the nodes on which the pods are created

```
1 kubectl get pods
2 kubectl get pods -o wide
```

Viewing POD Labels

```
1 kubectl get pods --show-labels
2 kubectl get pods --show-labels -n default
3 kubectl get pods --show-labels -n dev
```

Describing a POD

```
1 kubectl describe pod pod1 -n default
```

Finding POD Logs

```
1 kubectl logs pod1 -n default
```

Delete a pod

```
1 kubectl delete pod pod1 -n default
2 kubectl delete pod pod1 -n dev
```


Lab 5 Creating PODs using POD Object

Create a YAML file and deploy

```
1 nano pod.yaml
```

```
1  apiVersion: v1
2  kind: Pod
3  metadata:
4    name: mypod
5    labels:
6      app: myapp
7  spec:
8    containers:
9      - name: mycontainer
10      image: nginx:latest
11      ports:
12      - containerPort: 80
```

```
1 kubectl apply -f mypod.yaml
```

- `apiVersion` specifies the version of the Kubernetes API being used.
- `kind` specifies the type of Kubernetes resource, which in this case is a Pod.
- `metadata` contains information about the Pod, including its name and labels.
- `spec` specifies the desired state of the Pod, including the containers running in the Pod.
- `containers` contains a list of containers running in the Pod. Each container has a `name` and an `image` (the Docker image to use for the container).
- `ports` specifies the ports to expose from the container. In this case, port 80 is exposed.

How to know which API version to use

```
1 kubectl api-resources
```

List the pod status

```
1 kubectl get pods
```

Delete the pod

```
1 kubectl delete pod mypod -n default
```

Lab 6 Create PODs using ReplicationSet

Create a YAML and deploy

```
1 nano replicaset.yaml
```

```
1 apiVersion: apps/v1
2 kind: ReplicaSet
3 metadata:
4   name: my-replicaset
5 spec:
6   replicas: 3
7   selector:
8     matchLabels:
9       app: myapp
10  template:
11    metadata:
12      labels:
13        app: myapp
14    spec:
15      containers:
16      - name: mycontainer
17        image: nginx:latest
18        ports:
19      - containerPort: 80
```

```
1 kubectl apply -f replicaset.yaml
```

```
1 kubectl get pods
```

```
1 kubectl get rs
```

Delete ReplicaSet and PODs

```
1 kubectl delete rs my-replicaset
```

Lab 7 Create PODs using DeploymentSet

Create a YAML file and deploy

```
1 nano deployment.yaml
```

```
1 apiVersion: apps/v1
2 kind: Deployment
3 metadata:
4   name: my-deployment
5 spec:
6   replicas: 3
7   selector:
8     matchLabels:
9       app: myapp
10  template:
11    metadata:
12      labels:
13        app: myapp
14    spec:
15      containers:
16      - name: mycontainer
17        image: nginx:latest
18        ports:
19      - containerPort: 80
```

- `apiVersion` specifies the version of the Kubernetes API being used.
- `kind` specifies the type of Kubernetes resource, which in this case is a Deployment.
- `metadata` contains information about the Deployment, including its name.
- `spec` specifies the desired state of the Deployment, including the number of replicas to maintain.
- `replicas` specifies the number of desired replicas, which in this case is 3.
- `selector` specifies the labels used to select pods managed by the Deployment.
- `template` contains the pod template used to create new pods.
- `metadata` specifies the labels to apply to pods created by the Deployment.
- `spec` specifies the pod specification, including the container definition.
- `containers` contains a list of containers running in the pod. Each container has a name and an image.

```
1 kubectl apply -f deployment.yaml
```

```
1 kubectl get deployment
2 kubectl get pods
3 kubectl get pods -o wide
4 kubectl get rs
```

Increase replicas to 4

Edit the deployment.yaml file and make replicas to 4.

Reapply

```
1 kubectl apply -f deployment.yaml
2 kubectl get deployment
```

```
3 kubectl get rs
```

Delete Deployment

```
1 kubectl delete deployment my-deployment
```

Lab 8 Managing PODs

Deploy a pod

```
1 kubectl run pod1 --image=nginx --port=80 --labels=type=dev
```

Getting details about a particular POD

```
1 kubectl describe pod pod1
```

How to get the containers running on all PODs

```
1 kubectl get pods -o=custom-  
  columns=PodName:.metadata.name,Containers:.spec.containers[*].name,Image:.spec.containers[*].image
```

Attaching to a POD

```
1 kubectl get pods  
2 kubectl exec --stdin --tty <POD Name> -- /bin/bash
```

Attaching to a container in a POD

```
1 kubectl exec -i -t <pod name> --container <container name> -- /bin/bash
```

Create a pod on a particular node

Create a file called as `deploy-on-a-node.yml` and paste the definition

```
1 nano deploy-on-a-node.yml
```

```
1 apiVersion: v1  
2 kind: Pod  
3 metadata:  
4   name: pod1  
5 spec:  
6   containers:  
7   - name: container1  
8     image: nginx  
9   nodeName: node2
```

The **nodeName** field instructs the **Scheduler** to **assign the POD** to node2 computer.

Deploy the POD:

```
1 kubectl apply -f deploy-on-a-node.yml
```

Verify that the POD is created on the selected node:

```
1 kubectl get pods  
2 kubectl get pods -o wide
```

Add a label to a node

```
1 kubectl get nodes
2 kubectl get nodes -o wide
3 kubectl get nodes --show-labels
```

Choose one of your nodes, and add a label to it:

```
1 kubectl label nodes node2 disktype=ssd
```

Verify the label:

```
1 kubectl get nodes --show-labels
```

Filter all nodes that match your label:

```
1 kubectl get nodes --show-labels | grep disktype=ssd
```

Deploy a POD based on a label

We just added a label to node2: `disktype=ssd`

We want to deploy a POD to this node not by nodename but by the label.

Create a deployment file as shown below using nano:

```
1 nano deploy-with-label.yml
```

Paste the contents below:

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: pod1
5   labels:
6     env: dev
7 spec:
8   containers:
9   - name: container1
10     image: nginx
11   nodeSelector:
12     disktype: ssd
```

Perform the POD deployment:

```
1 kubectl apply -f deploy-with-label.yml
```

Verify the POD creation on node2, because it matches with the **NodeSelector** field `disktype` -

```
1 kubectl get pods -o wide
```

Remove a Label from a node

First find the node labels:

```
1 kubectl get nodes --show-labels
```

I want to delete the Label: `disktype`

```
1 kubectl label node <nodename> <labelName>-
```


Lab 9 Exposing PODs to Internet

Method-1 Using Load Balancer Service

Create a Deployment kind using this YAML

```
1 nano deployment.yaml
```

```
1 apiVersion: apps/v1
2 kind: Deployment
3 metadata:
4   name: website-deployment
5 spec:
6   replicas: 2
7   selector:
8     matchLabels:
9       app: webapp
10  template:
11    metadata:
12      labels:
13        app: webapp
14    spec:
15      containers:
16      - name: website-container
17        image: tanvisinghny/ssl-website
18        ports:
19        - containerPort: 80
20        - containerPort: 443
```

```
1 kubectl apply -f deployment.yaml
```

Verify if the Pods are created on both nodes

```
1 kubectl get pods -o wide
2 NAME                                READY   STATUS    RESTARTS   AGE   IP              NODE   NOMINATED
3 website-deployment-5fc6c46f9b-jw2xm 1/1     Running   0          40s   192.168.166.136 node1   <none>
4 website-deployment-5fc6c46f9b-s8s9c 1/1     Running   0          40s   192.168.104.21  node2   <none>
```

Create a Service object that exposes the deployment

```
1 kubectl expose deployment website-deployment --type=LoadBalancer --name=service-website-deployment
```

View the service

```
1 kubectl get services
2 NAME            TYPE        CLUSTER-IP   EXTERNAL-IP   PORT(S)    AGE
3 kubernetes       ClusterIP   10.96.0.1    <none>        443/TCP    3d
```



```
4 service-website-deployment LoadBalancer 10.99.3.87 <pending> 80:30923/TCP,443:30274/TCP 20s
```

Note that External IP is pending

Patch the service with node1 and node2 host IPs

```
1 kubectl patch svc service-website-deployment -n default -p '{"spec": {"type": "LoadBalancer", "externalIPs": ["10.0.0.5", "10.0.0.6"]}}'
```

View the service again

```
1 kubectl get services
2 NAME                                TYPE                CLUSTER-IP    EXTERNAL-IP    PORT(S)                AGE
3 kubernetes                          ClusterIP           10.96.0.1     <none>         443/TCP                3d
4 service-website-deployment          LoadBalancer        10.99.3.87    10.0.0.5,10.0.0.6 80:30923/TCP,443:30274/TCP 3m21s
```

On the VMs open HTTP, HTTPS Ports and access the website using the node1 and node2 IP addresses



Method-2 Using NodePort Service

Delete the Service you created above

```
1 kubectl delete service service-website-deployment
```

View if the service is deleted

```
1 kubectl get services
2 NAME                                TYPE                CLUSTER-IP    EXTERNAL-IP    PORT(S)                AGE
3 kubernetes                          ClusterIP           10.96.0.1     <none>         443/TCP                3d
```

Create a Deployment kind using this YAML

```
1 nano deployment.yaml
```

```
1 apiVersion: apps/v1
2 kind: Deployment
3 metadata:
4   name: website-deployment
5 spec:
6   replicas: 2
7   selector:
8     matchLabels:
9       app: webapp
10  template:
11    metadata:
12      labels:
13        app: webapp
14    spec:
15      containers:
16      - name: website-container
17        image: tanvisinghny/ssl-website
18        ports:
19        - containerPort: 80
20        - containerPort: 443
21  ---
22 apiVersion: v1
23 kind: Service
24 metadata:
25   name: website-service
26 spec:
27   type: NodePort
28   selector:
29     app: webapp
30   ports:
31   - port: 80
32     targetPort: 80
33     nodePort: 30080
34     protocol: TCP
35     name: http
36   - port: 443
37     targetPort: 443
38     nodePort: 30443
39     protocol: TCP
40     name: https
41   externalIPs:
42   - 10.0.0.5
43   - 10.0.0.6
```

```
1 kubectl apply -f deployment.yaml
```

View the created objects

```
1 kubectl get deployments
2 NAME                READY  UP-TO-DATE  AVAILABLE  AGE
3 website-deployment  2/2    2            2           14s
4
5 kubectl get service
```

6	NAME	TYPE	CLUSTER-IP	EXTERNAL-IP		PORT(S)		AGE	
7	kubernetes	ClusterIP	10.96.0.1	<none>		443/TCP		3d1h	
8	website-service	NodePort	10.97.136.212	10.0.0.5,10.0.0.6		80:30080/TCP,443:30443/TCP		21s	
9									
10	kubectl get pods -o wide								
11	NAME		READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED
	NODE	READINESS GATES							
12	website-deployment-5fc6c46f9b-hfgzp		1/1	Running	0	81s	192.168.104.24	node2	<none>
	<none>								
13	website-deployment-5fc6c46f9b-lb59l		1/1	Running	0	81s	192.168.166.139	node1	<none>
	<none>								

Access the website in the browser



Cleanup

```
1 kubectl delete deployment website-deployment
2 kubectl delete service website-service
```

Method-3 Using ClusterIP

Create a Deployment kind using this YAML

```
1 nano deployment.yaml
```

```
1 apiVersion: apps/v1
2 kind: Deployment
3 metadata:
4   name: website-deployment
5 spec:
6   replicas: 2
7   selector:
```

```

8     matchLabels:
9       app: webapp
10  template:
11    metadata:
12      labels:
13        app: webapp
14    spec:
15      containers:
16      - name: website-container
17        image: tanvisinghny/ssl-website
18        ports:
19        - containerPort: 80
20        - containerPort: 443
21  ---
22  apiVersion: v1
23  kind: Service
24  metadata:
25    name: website-service
26  spec:
27    type: ClusterIP
28    selector:
29      app: webapp
30    ports:
31    - port: 80
32      protocol: TCP
33      name: http
34    - port: 443
35      protocol: TCP
36      name: https
37    externalIPs:
38    - 10.0.0.5
39    - 10.0.0.6

```

```
1 kubectl apply -f deployment.yaml
```

View the created objects

```

1 kubectl get deployments
2 NAME                READY   UP-TO-DATE   AVAILABLE   AGE
3 website-deployment  2/2     2             2           69s
4
5 kubectl get service
6 NAME                TYPE        CLUSTER-IP   EXTERNAL-IP   PORT(S)          AGE
7 kubernetes          ClusterIP   10.96.0.1    <none>        443/TCP          3d2h
8 website-service     ClusterIP   10.99.86.166 10.0.0.5,10.0.0.6 80/TCP,443/TCP   16s

```

Access the website in the browser



Lab 10 Kubernetes Networking

Finding POD IP Range

```
1 kubectl get ipamblocks.crd.projectcalico.org
```

For example

```
1 kubectl get ipamblocks.crd.projectcalico.org
2 NAME                                AGE
3 192-168-104-0-26                    82m
4 192-168-166-128-26                 83m
5 192-168-219-64-26                  101m
```

Default CIDR Block is /26 = 64 IPs

This allows to create 64 PODS only

Changing Block CIDR

Step-1 Install calicoctl as a Kubernetes pod

```
1 kubectl apply -f https://docs.projectcalico.org/manifests/calicoctl.yaml
```

Verify

```
1 kubectl get pods -n kube-system
2 NAME                                READY   STATUS    RESTARTS   AGE
3 calico-kube-controllers-658d97c59c-gzqz4 1/1     Running   0          111m
4 calico-node-9dnxq                        1/1     Running   0          90m
5 calico-node-jc25j                        1/1     Running   1 (96m ago) 111m
6 calico-node-k5l67                        1/1     Running   0          91m
7 calicoctl                               1/1     Running   0          19s
8 coredns-76f75df574-44bz7                1/1     Running   0          136m
9 coredns-76f75df574-44g76                1/1     Running   0          136m
10 etcd-master                             1/1     Running   1 (96m ago) 136m
11 kube-apiserver-master                    1/1     Running   1 (96m ago) 136m
12 kube-controller-manager-master           1/1     Running   3 (96m ago) 136m
13 kube-proxy-4fwg6                         1/1     Running   0          90m
14 kube-proxy-64sfl                         1/1     Running   0          91m
15 kube-proxy-w68zd                         1/1     Running   1 (96m ago) 136m
16 kube-scheduler-master                    1/1     Running   3 (96m ago) 136m
```

Set an alias:

```
1 alias calicoctl="kubectl exec -i -n kube-system calicoctl -- /calicoctl "
```

Step-2 Add a new IP pool

```
1 calicoctl create -f -<<EOF
2 apiVersion: projectcalico.org/v3
3 kind: IPPool
4 metadata:
5   name: new-pool
6 spec:
7   cidr: 10.0.0.0/8
```

```
8   ipipMode: Always
9   natOutgoing: true
10 EOF
```

Lab 12 Taints and Tolerations

POD scheduling issues

- Pod scheduling is one of the most important aspects of Kubernetes administration.
- Effective scheduling can improve performance, reduce costs, and make clusters easier to manage.
- Taints and tolerations help prevent your pods from scheduling to undesirable nodes.
- When we create a pod, the scheduler in the control plane looks at the available nodes and verifies conditions before assigning a pod to the nodes.
- If there are no errors during the verification, the pod will be scheduled on the node.
- If the conditions of the verification aren't satisfied, then your pods will be put in a **Pending** state.
- You can use `kubectl describe pods <pod-Name>` and scroll down to Events to find the precise reasons for the pending state
- You can run `kubectl get events`.
- If it is a hardware resource issue, you may get some error about CPU or Memory or Storage.

Taint

Taints are a Kubernetes node property that enable nodes to repel certain pods.

Master Node Example:

By default except Kubernetes System PODs, User PODs are not created on Control-plane node.

This because the master node is has a Taint:

```
node-role.kubernetes.io/control-plane:NoSchedule
```

The NoSchedule Taint REPELS the Scheduler and does not allow POD creation.

Since Taints is a NODE Property, you can see a node taints using describe command:

```
1 kubectl describe node master
```

The **taint effect** defines how a tainted node reacts to a pod without appropriate toleration. It must be one of the following effects;

- **NoSchedule** — The pod will not get scheduled to the node without a matching toleration.
- **NoExecute** — This will immediately evict all the pods without the matching toleration from the node.
- **PerferNoSchedule** — This is a softer version of NoSchedule where the controller will not try to schedule a pod with the tainted node. However, it is not a strict requirement.

Toleration

- **Toleration** is applied to pods and allows the pods to schedule onto nodes with matching taints.

POD (Toleration) ===Scheduled===> NODE (Taint)

For the above control-plane example, to schedule PODs on control-plane node add a Toleration:

```
1 tolerations:
2   - key: "node-role.kubernetes.io/control-plane"
3     effect: "NoSchedule"
4
```

Actual script:


```
1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    name: deployment1
5  spec:
6    replicas: 20
7    selector:
8      matchLabels:
9        app: app1
10   template:
11     metadata:
12       labels:
13         app: app1
14     spec:
15       containers:
16       - name: app1-con
17         image: nginx:latest
18         ports:
19         - containerPort: 80
20       tolerations:
21       - key: "node-role.kubernetes.io/control-plane"
22         effect: "NoSchedule"
```

Taint a Node

```
1  kubectl taint nodes node1 gpu=true:NoSchedule
```

We want to ensure only Deployments with Tolerations matching `gpu=true:NoSchedule` should allow scheduler to deploy PODs on node1.

Try this script

```
1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    name: deployment1
5  spec:
6    replicas: 20
7    selector:
8      matchLabels:
9        app: app1
10   template:
11     metadata:
12       labels:
13         app: app1
14     spec:
15       containers:
16       - name: app1-con
17         image: nginx:latest
18         ports:
19         - containerPort: 80
20       tolerations:
21       - key: "gpu"
22         operator: "Equal"
23         value: "true"
24         effect: "NoSchedule"
```

If that doesn't work, change Tolerations as:

Lab 15 Kubernetes Backup

Download velero

```
1 curl -LO https://github.com/vmware-tanzu/velero/releases/download/v1.10.1/velero-v1.10.1-linux-amd64.tar.gz
```

Extract

```
1 tar zxvf velero-v1.10.1-linux-amd64.tar.gz
```

Move to bin path

```
1 sudo mv velero-v1.10.1-linux-amd64/velero /usr/local/bin/velero
```

Backup a namespace and its objects

```
1 velero backup create <backup-name> --include-namespaces <namespace>
```

Backup all deployments in the cluster

```
1 velero backup create <backup-name> --include-resources deployments
```

Backup the deployments in a namespace

```
1 velero backup create <backup-name> --include-resources deployments --include-namespaces <namespace>
```

Backup entire cluster including cluster-scoped resources

```
1 velero backup create <backup-name>
```

Backup a namespace and include cluster-scoped resources

```
1 velero backup create <backup-name> --include-namespaces <namespace> --include-cluster-resources=true
```

Include resources matching the label selector

```
1 velero backup create <backup-name> --selector <key>=<value>
```

Include resources that are not matching the selector

```
1 velero backup create <backup-name> --selector "<key> notin (<value>)"
```

Exclude kube-system from the cluster backup

```
1 velero backup create <backup-name> --exclude-namespaces kube-system
```

Exclude secrets from the backup

```
1 velero backup create <backup-name> --exclude-resources secrets
```

Exclude secrets and rolebindings

```
1 velero backup create <backup-name> --exclude-resources secrets,rolebindings
```

Restore two namespaces and their objects

```
1 velero restore create <backup-name> --include-namespaces <namespace1>,<namespace2>
```

Restore all deployments and configmaps in the cluster

```
1 velero restore create <backup-name> --include-resources deployments,configmaps
```

Restore only namespaced resources in the cluster

```
1 velero restore create <backup-name> --include-cluster-resources=false
```